

Latent Outlet Potential Estimation

Data Storm 7.0 | Storming Round

Sri Lanka | January 2026 horizon

Submitted by Smile Labs

Powered by OCTAVE – John Keells Group | Rotaract Moratuwa

One-line method

We estimate $\mathbb{E}[\text{true_demand}_i \mid X_i]$ where observed sales $y_i = \min(\text{true_demand}_i, \text{constraint}_i)$ are right-censored. The estimand is *not point-identified* from observational data alone [1]; we therefore report [lower, point, upper] per outlet, with the point combining Stochastic Frontier Analysis [2, 3], multi-quantile gradient boosted regression, censoring correction [4], Conformalised QR [5], and bootstrap-derived peer-bucket caps.

Headline	Value
Outlets predicted (full population in scope)	20,000
External POIs scraped (Geofabrik PBF, 9 categories)	42,386
Outlets with ≥ 1 POI within 2 km	80.9 %
Rejected check records (coords + txn + holidays)	10,179 across 3 datasets
Median uplift vs. historical max	1.250×
Submission preflight (6 automated checks)	6 / 6 PASS
Method stack methods	SFA + multi-q XGBoost + CH-3 + CQR + Manski

Build trail

The v2 stack was built and audited across **seven rounds of AI council reviews** (four parallel premium-model critics per round). Provenance: `Reviews/council_review.md` (R1), `council_round{2..7}/` (R2-R7). Every fix in `src/` cites the council finding it addresses (FIX_R7, FIX_R4, FIX_M1, etc.).

1 Data Forensics and Hygiene

Bronze → Silver → Gold lakehouse

Raw CSVs ingest *as-is* into data/bronze/ with a SHA-256 audit log. The Silver layer applies six **reusable, parameterizable** DQ check functions (src/quality/checks.py) consistently across all five raw datasets; failed rows are quarantined to data/silver_rejected/ with dataset_name, failed_check, and failure_reason. Gold is the model-ready outlet × feature parquet.




Figure 1: Bronze → Silver → Gold pipeline with rejected-records store.

Reusable DQ check functions applied to every dataset

Function (checks.py)	What it catches
duplicate_check(df, key_columns)	duplicate composite keys
null_check(df, columns)	NaN or empty mandatory fields
range_check(df, column, min, max)	numeric out-of-range
domain_check(df, column, allowed_values)	misspellings, unexpected enums
referential_integrity_check(df, column, ref)	foreign keys not in master
geospatial_bounds_check(df, lat, lon, lat_range, lon_range)	coords outside Sri Lanka

Table 1: All six functions are parameterised; zero check is hardcoded for one dataset.

Legacy SFA / ERP system artifacts neutralised

Artifact	Records	Action
Outlet_Type typos (Grocery, Bakery) + Outlet_Size lowercase + missing	1,581	normalised / mapped to Unknown
Outlet coordinates outside Sri Lanka bounds (5.5–10°N, 79–82.5°E)	240	quarantined
Transactions with non-positive Volume_Liters / Total_Bill_Value	9,606	quarantined (some rows fail both checks)
Holiday duplicate (date + name + type) rows	93	deduped, originals to rejected

Table 2: Every quarantined row carries a documented failure_reason.

Capacity feature engineered from a measured saturation knee

Volume *stops growing* after 3 coolers per outlet (notebook 23): median monthly volume at 3, 4, and 5+ coolers is statistically indistinguishable. We therefore use Cooler_Count as the deterministic sign anchor for the PCA capacity component (src/modeling/constraint_score.py).




Figure 2: Median monthly volume vs. Cooler_Count: knee at 3 anchors the capacity feature.

2 POI Data Acquisition

Geofabrik PBF » naïve Overpass

The official PDF explicitly judges “*how robust was the web-scraping / API pipeline for external POI data*”. Naïvely querying Overpass for $20,000 \times 9 \text{ cats} \times 4 \text{ radii} = 720,000$ requests is infeasible (timeouts + ToS violation). We instead download the **Geofabrik Sri Lanka country extract** [6] (`sri-lanka-latest.osm.pbf`, ~136 MB) once, parse it locally with `pyrosm`, and run nearest-neighbour searches with `sklearn.neighbors.BallTree(metric="haversine")`. End-to-end: 25–45 min, entirely offline after the single fetch. The pipeline lives in `poi_pipeline/` as a separate, idempotent sub-project.

Nine target categories + 80.9 % outlet coverage

Category	POIs	OSM tags
Schools / universities	5,694	amenity in {school, university, college, kindergarten}
Transport hubs	5,799	highway=bus_stop, amenity=bus_station, railway=station
Hospitals + healthcare	2,007	amenity in {hospital, clinic, pharmacy}, healthcare=*
Restaurants / cafes	4,667	amenity in {restaurant, cafe, fast_food, food_court}
Supermarkets / groceries	2,543	shop in {supermarket, convenience, grocery}, amenity=marketplace
Religious places	9,255	amenity=place_of_worship
Hotels / attractions	5,513	tourism in {hotel, guest_house, attraction}
Offices	4,169	office=*, amenity in {townhall, post_office}
Banks / ATMs	2,739	amenity in {bank, atm}
Total	42,386	80.9 % of outlets with ≥ 1 POI within 2 km

Table 3: POI extraction summary from `poi_pipeline/data/pois/`.

Honest coverage + signal disclosure (cited in viva)

OSM-LK mapping bias. Well-mapped: schools, hospitals, banks, hotels, religious places. Patchy: small *kades*, independent grocers, informal roadside stalls – exactly the outlet types we predict for.

Raw POI signal is weak per outlet. Our own EDA (notebook 23) found POI counts correlate weakly with monthly volume ($|r| \leq 0.026$ across all Outlet_Type groups, Figure 3). POI therefore enters the model only via the composite `poi_catchment_score` (z-averaged 9-category Gaussian-decay scores) at ~20 % weight in the constraint score – not as a standalone demand predictor.




Figure 3: POI $\times$ Outlet_Type interaction (notebook 23). Per-outlet POI counts have low absolute correlation with monthly volume regardless of channel, confirming POI’s role as a *catchment* signal rather than a direct predictor.

3 Causal / Probabilistic Methodology

Identification analysis (right-censored, not point-identified)

The estimand $\mathbb{E}[\text{true_demand}_i \mid X_i]$ is **not point-identified** from observational data alone [1]. Observed $y_i = \min(\text{true_demand}_i, \text{constraint}_i)$ is right-censored: we only see a lower bound for any outlet that hit a constraint. We embrace this honestly and report a *band*, not a single number, alongside the point estimate.




Figure 4: DAG. Solid arrows are observed; dashed arrow indicates partial identification of the estimand from $\min(\cdot)$.

Figure 5 shows the empirical evidence motivating the $\max(\cdot, \text{observed_max})$ floor in our final formula: for 15.62 % of outlets, the q90 frontier estimate lies *below* the outlet's observed historical maximum – without the floor, V3b validation fails for those rows.




Figure 5: Distribution of $\text{frontier_q90} - \text{observed_max_monthly_liters}$. Negative values (15.62 % of outlets, worst at Extra Large at 32.24 %) motivate the floor at observed_max in our formula.

Method stack (four components, three boxes)

Robust lower bound + frontier ensemble

$\text{lower_bound}_i = \max(3\text{rd-highest month}_i, \text{median}_i)$. **Frontier** = 60/40 blend of (a) **XGBoost 2.0** [7] with `reg:quantileerror, quantile_alpha=[0.5, 0.75, 0.9, 0.95]`, monotone constraints + isotonic post-sort [8], and (b) **SFA** [2]: $y = X\beta + v - u$ with truncated-normal u , direct MLE; per-outlet TE = $\exp(-\mathbb{E}[u \mid \varepsilon])$ [3] (fitted $\sigma_v \approx 0.41$, $\sigma_u \approx 0.78$, $\lambda \approx 1.92$; median TE ≈ 0.84). Leaky observed_ aggregates dropped from SFA design matrix.

Censoring correction + constraint score

Chernozhukov-Hong (2002) 3-step [4] (extending Powell's censored quantile regression [9]): propensity $P(\text{censored} \mid X)$ fit on plateau / stuck-at-ceiling proxies; q90 refit on $P < 0.10$. Note: only 1.16 % of outlets show a true plateau fingerprint globally (primarily `DIST_S_01, DIST_S_02`); CH-3 is a small-population correction, not a load-bearing transform. **Constraint score** $s_i \in [0, 1]$ from three orthogonal signals: peer-q90 frontier residual z-score, plateau gate, PCA-decorrelated capacity. **Outlets with $s_i \geq 0.40$ receive a 1.25× uplift floor over observed_max** , bounded by the bucket cap.

Final formula + calibrated interval

$$\text{potential}_i = \max(\text{obs_max}_i, \text{lower}_i + s_i(F_i - \text{lower}_i), \mathbf{1}_{[s_i \geq 0.4]} \cdot 1.25 \cdot \text{obs_max}_i), \text{ capped at } B_i \cdot \text{obs_max}_i$$

F_i is the SFA + multi-q XGBoost frontier; s_i is the constraint score; B_i is the empirical 95th-pct uplift inside outlet i 's (Outlet_Type × Outlet_Size) bucket. We additionally report a calibrated $[q_{05}, q_{95}]$ via Conformalised QR [5] on a 20 % outlet-level holdout (disjoint from the q90 training set after the R4 fix), plus Manski worst-case bounds [1] per outlet.

4 Validation, Outputs, and GenAI Transparency

Validation without ground truth (6-item auto-checklist)

#	Check	Detail (final run)	Status
V1	schema + row count	[Outlet_ID, Maximum_Monthly_Liters], 20,000 rows	OK
V2	no NaN, no negatives, unique IDs	NaN=0, neg=0, unique=True	OK
V3a	every Outlet_ID exists in outlet_master	missing = 0	OK
V3b	predicted ≥ historical max for ≥ 99 % of outlets	0.00 % below historical max	OK
V4	median uplift in [1.25, 2.2]	median = 1.250	OK
V5	cap-binding rate < 25 % (bucket-specific cap)	0.00 % at cap	OK

Table 4: Live output of src/reporting/validation.py (Results/validation_report.md).

Manski band coverage + sensitivity




Figure 6: Per-outlet Manski band. **Sensitivity sweep:** across quantile {0.85, 0.9, 0.95} × scheme {balanced, frontier-heavy, plateau-heavy} × cap {2..6}, median uplift stays at 1.250× and cap-binding at 0 %.

Generative AI workflow (20 % rubric)

Phase	AI usage	Human validation applied
Research swarm	10-channel parallel research swarm (Claude Opus 4 + GPT-5.5) producing research/research_brief.md	each claim cross-checked against cited URLs
Methodology audit	Seven rounds of 4–5 parallel critics (Statistician, Skeptic, Architect, Safety+DE, Gap, Visual, Judge, Risk)	every finding cites file:line; verified or rebutted manually
Code + design	Cursor routing scaffolded src/; LLM-assisted JLMS closed form; OSM tag survey; Geofabrik vs Overpass trade-off	syntax + import check; ReadLints clean; SFA cross-checked against Greene; pipeline timed end-to-end
Report drafting	LaTeX template assembly + content port from markdown source	every claim traced to an artifact in Results/ or data/gold/

Table 5: Per-phase GenAI transparency. Full log: Docs/ai_transparency_log.md (R1-R7 audit trail).

Deliverables

Artifact	Path
Platform submission CSV	Results/smile_labs_predictions.csv (20,000 rows, [Outlet_ID, Maximum_Monthly_Liters])
Full business output	Results/smile_labs_predictions_full_20000.csv
Manski + conformal bands	Results/manski_bands_v2.csv, conformal_intervals_v2.csv
Validation	Results/validation_report.md (6/6 PASS) + validation_extended_v6.md (V4b/V6/V7)
Reproducible codebase	run_pipeline.py one-shot, or Notebooks/20..22_v2_*.ipynb; modules in src/

[1] C. F. Manski. *Partial Identification of Probability Distributions*. Springer, 2003. DOI: 10.1007/b97478.

[2] D. Aigner, C. A. K. Lovell, and P. Schmidt. “Formulation and Estimation of Stochastic Frontier Production Function Models”. In: *Journal of Econometrics* 6.1 (1977), pp. 21–37. DOI: 10.1016/0304-4076(77)90052-5.

[3] J. Jondrow, C. A. K. Lovell, I. S. Materov, and P. Schmidt. “On the Estimation of Technical Inefficiency in the Stochastic Frontier Production Function Model”. In: *Journal of Econometrics* 19.2–3 (1982), pp. 233–238. DOI: 10.1016/0304-4076(82)90004-5.

[4] V. Chernozhukov and H. Hong. “Three-Step Censored Quantile Regression and Extramarital Affairs”. In: *Journal of the American*

Statistical Association 97.459 (2002), pp. 872–882. DOI: 10.1198/016214502388618663.

[5] Y. Romano, E. Patterson, and E. J. Candès. “Conformalized Quantile Regression”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. Vol. 32. 2019. URL: https://papers.nips.cc/paper_files/paper/2019/hash/5103c3584b063c431bd1268e9b5e76fb-Abstract.html.

[6] Geofabrik GmbH. *Geofabrik OpenStreetMap Data Extracts – Sri Lanka*. Accessed January 2026 via Sri Lanka country extract (sri-lanka-latest.osm.pbf, 136 MB). 2026. URL: <https://download.geofabrik.de/asia/sri-lanka.html>.

[7] Tianqi Chen and Carlos Guestrin. “XGBoost: A Scalable Tree Boosting System”. In: *Proc. 22nd ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining (KDD ’16)*. ACM, 2016, pp. 785–794. DOI: 10.1145/2939672.2939785.

[8] V. Chernozhukov, I. Fernández-Val, and A. Galichon. “Quantile and Probability Curves Without Crossing”. In: *Econometrica* 78.3 (2010), pp. 1093–1125. DOI: 10.3982/ECTA7880.

[9] J. L. Powell. “Censored Regression Quantiles”. In: *Journal of Econometrics* 32.1 (1986), pp. 143–155. DOI: 10.1016/0304-4076(86)90016-3.